

FORMATTING INSTRUCTIONS FOR ICLR 2027

CONFERENCE SUBMISSIONS

Anonymous authors

Paper under double-blind review

ABSTRACT

1 INTRODUCTION

2 RELATED WORKS

3 PRELIMINARIES

3.1 MULTI-AGENT COORDINATION AS BELIEF-CONDITIONED SEQUENTIAL DECISION MAKING

In sequential multi-agent decision making, the consequence of an ego action depends not only on the physical state but also on how other agents will respond to that action. This dependence is especially explicit when their goals, incentives, or other behaviorally relevant information are not directly observable. We therefore adopt the interactive POMDP (I-POMDP) perspective, which extends sequential decision making with models of other agents and maps beliefs over interactive states to actions (Gmytrasiewicz & Doshi, 2005). This is a well-established modeling structure: model the partner, predict its behavior, choose an action under that prediction, and update the model from subsequent interaction (Carmel & Markovitch, 1996). Recent work also provides an instance of the same overall coupling by combining inferred partner goals with goal-conditioned policies and planning (Huang et al., 2024).

Let s_t denote the public interaction state at decision time t , including the physical state, public history, turn phase, and any pending interaction. Agent i receives observation o_t^i and has a private latent state θ_i ; for a partner j , θ_j denotes the task-relevant component of its latent model, such as a hidden goal, incentive, or other information affecting future behavior. We then use this perspective to separate two research objects. First, a *belief* is the player’s information state about the feasible or plausible latent states of other players. Given public evidence e_t^i , the belief evolves through the standard model-update interface

$$b_{t+1}^i = \text{Update}(b_t^i, e_t^i; s_t, a_t^i), \quad (1)$$

where e_t^i may include an observed partner action and its public state consequence.

Second, *conditioned planning* evaluates candidate ego actions using the partner model induced by the current belief. For a legal action $a \in \mathcal{A}_i(s_t)$, let $\hat{\pi}_j^i(y | s_t, a, b_t^i)$ denote the resulting predictive model for partner action or response y . The player then selects an action by optimizing its expected continuation utility:

$$a_t^i \in \arg \max_{a \in \mathcal{A}_i(s_t)} Q_i(s_t, b_t^i, a), \quad Q_i(s_t, b_t^i, a) = \mathbb{E}[U_i(s_T; \theta_i) | s_t, o_t^i, b_t^i, a]. \quad (2)$$

The expectation includes the predicted partner responses and subsequent interaction dynamics. The two objects are therefore distinct but coupled: b_t^i determines the partner model used by planning, while the realized interaction outcome supplies evidence for the next belief.

3.2 BENAC-BASED ENVIRONMENT AND ORACLE PARTNERS

Base BENAC. We build on BENAC, a negotiation setting in which players make sequential, action-level commitments toward multiple goals (Benac et al., 2026). The original implementation stores binding binary commitments in a policy matrix, represents goals by satisfaction masks, applies an accepted bilateral deal by a bitwise OR, and advances the schedule without changing commitments when a deal is rejected. The extension used here preserves this commitment semantics while making private incentives, explicit responses, menu offers, and multi-player scheduling part of the game specification.

Hidden incentives. An instance is $\mathcal{M} = (N, \mathbf{K}, \mathcal{G}, V, \rho, M, F)$, where $N \geq 2$ is the number of players, $\mathbf{K} = (K_1, \dots, K_N)$ gives the number of binary actions per player, $\mathcal{G} = \{0, \dots, G-1\}$ is the set of public goals, ρ is the public proposer schedule, M is the per-side commitment-change budget, and F is an optional set of forbidden actions. The environment stores the complete preference matrix

$$V = (v_{i,g}) \in \{-1, 0, +1\}^{N \times G}, \quad +1 = \text{WANT}, \quad 0 = \text{NEUTRAL}, \quad -1 = \text{AVOID}. \quad (3)$$

Player i privately observes only its own row $v_{i,:}$ in the normal private view. For each player i , let $c_{t,i} \in \{0, 1\}^{K_i}$ be its binding commitments and let $C_t = (c_{t,1}, \dots, c_{t,N})$. A public state is $s_t = (C_t, t, \tau_t)$ together with the public game specification, where τ_t is the transcript. Commitments are monotone: every transition preserves existing 1 entries and can only add 0→1 entries. Each goal g has a nonempty public requirement set $\Gamma_g \subseteq \{(i, k) : 1 \leq i \leq N, 1 \leq k \leq K_i\}$ and is satisfied iff

$$z_g(C_t) = \mathbf{1}[c_{t,i,k} = 1 \text{ for all } (i, k) \in \Gamma_g]. \quad (4)$$

The implementation restricts every goal to involve at least two distinct players. Hence, in the current partial-information setting, the hidden partner state is the preference/incentive row rather than the goal graph: goal requirements remain public. Player i 's terminal utility is the additive preference-weighted utility

$$U_i(C_T) = \sum_{g=0}^{G-1} v_{i,g} z_g(C_T). \quad (5)$$

The episode terminates exactly when the schedule is exhausted, $t = T = |\rho|$; goal completion does not terminate the episode early. The runner reports the utility at this terminal state and does not emit a separate shaped transition reward.

Menu offers. At turn t , the proposer is $p = \rho_t$. An ordinary offer addressed to $q \neq p$ is

$$o = (q, \bar{c}_p, \bar{c}_q), \quad \bar{c}_p \in \{0, 1\}^{K_p}, \quad \bar{c}_q \in \{0, 1\}^{K_q}, \quad (6)$$

where the vectors are final target vectors rather than deltas. If $D_k = \{a : c_{t,k,a} = 0, \bar{c}_{k,a} = 1\}$, legality requires preservation of existing commitments, $|D_p| \leq M$, $|D_q| \leq M$, no newly added forbidden action, and $|D_p| + |D_q| > 0$. The ordinary action space contains PASS and all legal offers to every $q \neq p$. PASS leaves C_t unchanged and consumes the turn. For an ordinary offer, the partner response is $r \in \{\text{ACCEPT}, \text{REJECT}\}$. ACCEPT updates

$$c_{t+1,k} = c_{t,k} \vee \bar{c}_k \quad (k \in \{p, q\}), \quad (7)$$

leaving all other rows unchanged; REJECT leaves C_t unchanged. Both outcomes advance the schedule cursor by exactly one and are recorded in the public transcript.

When menus are enabled, a menu is a pair of distinct legal ordinary offers to the same partner, $m = (o^{(1)}, o^{(2)})$. Its response set is $\{\text{REJECT}, \text{CHOOSE_1}, \text{CHOOSE_2}\}$; only the chosen option binds, rejection binds neither, and the menu also consumes one proposer turn.

Multi-party interaction and observation. The episode is multi-party even though each individual offer is bilateral: N may exceed two, a goal may require commitments from multiple players, and commitments accumulated in different bilateral interactions jointly determine the terminal utilities. The generator defaults to $(N, K_i, G, R, M) = (3, 3, 8, 4, 1)$, samples a player permutation for each of R rounds, and uses the concatenated schedule ρ ; all of these dimensions and the schedule are configurable. The schema permits $N = 2$ for small games, while the generator directly supports $N \geq 3$. Each resolved proposal consumes one scheduled turn, with a response only when the proposal is an offer.

The normal private observation of player i contains the public goals and requirements, binding commitments, schedule and cursor, public transcript, and, at the relevant phase, legal partners/offers or a pending offer, together with $v_{i,:}$. A pending offer is hypothetical until its response is chosen. The public observation omits private preferences; an explicit full view exposes all rows only for debugging. The current implementation has no free-form message channel or explicit reveal action, so its public evidence consists of structured proposals, responses, and the resulting commitment snapshot.

Belief solver. For ego i , the finite latent type space for partner j is the preference-row space $\Theta_j = \{\text{WANT}, \text{NEUTRAL}, \text{AVOID}\}^G$. The solver maintains joint hypotheses

$$\mathcal{H}_t^i \subseteq \prod_{j \neq i} \Theta_j, \quad \mathcal{B}_t^i = (\mathcal{H}_t^i, w_t^i), \quad \mathcal{H}_{j,t}^i = \{h_j : \exists h \in \mathcal{H}_t^i\}, \quad (8)$$

where $\mathcal{H}_{j,t}^i$ is the currently feasible set for partner j and the joint representation preserves cross-partner correlations. The concrete `BeliefState` stores normalized log weights. The default `ConditionalPreferencePrior` starts from label probabilities $(0.4, 0.2, 0.4)$ in `WANT/NEUTRAL/AVOID` order and conditions on the ego’s row and the generator’s validity constraints; controlled fixtures may instead provide a finite support explicitly.

Upon a valid partner-generated proposal or response, the exact filter scores each hypothesis with the known partner kernel and applies

$$w_{t+1}^i(h) = \frac{w_t^i(h) \ell_j(e_t | s_t, \theta_j(h))}{\sum_{h' \in \mathcal{H}_t^i} w_t^i(h') \ell_j(e_t | s_t, \theta_j(h'))}, \quad h \in \mathcal{H}_t^i, \quad (9)$$

with zero-likelihood hypotheses assigned zero posterior mass. The concrete implementation retains such hypotheses in the enumerated support with $-\infty$ log weight. Ego actions are interventions and do not update the likelihood. The filter replays only unseen, valid, in-order public events and pending offers; invalid or fallback events are marked as such and are not evidence. In the repository, the kernel is the known stochastic `SoftProgressPolicy`, which uses the acting player’s own preference row and public state, scores actions by accepted state potential gain (with partial-progress weight 0.5, zero on the last turn), and applies a temperature-0.5 softmax with a 0.02 uniform mixture. It is a predictive partner kernel, not a rational oracle and not a model of a partner’s belief about other players. The deterministic commitment snapshot is used to validate the public transition; the likelihood is attached to the observed proposal or response. The associated `BayesPlanner` consumes $\mathcal{B}_t^i$, branches on the same partner kernel, updates the belief at non-ego nodes, and maximizes ego terminal utility over the remaining schedule without discount; it is therefore a belief-conditioned reference planner rather than an equilibrium solver.

Oracle partners. An observation-limited rational partner would have the form

$$a_j^t = \pi_{\text{oracle}}(o_j^t, \mathcal{B}_{-j}^t), \quad (10)$$

where o_j^t contains the public state and j ’s own private incentive, and $\mathcal{B}_{-j}^t$ is maintained from information available to j when inferring other players. Under this semantics, ground-truth access is limited to j ’s own private state; all inference about other agents must use the same observation and belief-update interface as an ordinary player.

The current reference implementation should be distinguished from this observation-limited definition. `RationalOraclePolicy` is bound by `GameRunner` to the live state and player identity, then delegates to `PerfectInfoSolver`. The solver reads the complete `GameSpec`, including every row of the ground-truth V , and performs finite-horizon backward induction. For a pending ordinary offer o to responder q , let $W_q^{\text{acc}}(s, o)$ and $W_q^{\text{rej}}(s, o)$ denote the continuation utilities obtained by accepting and rejecting. Its response is

$$r_q^{\text{PI}}(s, o) = \begin{cases} \text{ACCEPT}, & W_q^{\text{acc}}(s, o) \geq W_q^{\text{rej}}(s, o), \\ \text{REJECT}, & \text{otherwise.} \end{cases} \quad (11)$$

At proposer nodes it chooses an accepted offer maximizing the proposer’s continuation utility; PASS is the baseline and wins ties, with remaining ties resolved by legal enumeration order. This solver is menu-disabled. It does not maintain a belief about the ego, and the observation passed to the policy is used only for identity and turn-consistency checks; legality is checked against the bound live state rather than used to restrict the solver’s access to V .

For clarity, the implementation exposes four distinct objects: (i) the *environment ground truth*, comprising the full V , public specification, and current physical state; (ii) the *agent-observable information*, comprising the public state plus the agent’s own private row in private mode; (iii) the *solver-maintained belief*, comprising the current feasible partner hypotheses and weights; and (iv) the *policy*, which maps the relevant information to an action. The exact filter and `BayesPlanner` use the belief object, whereas the current `RationalOraclePolicy` is a full-information reference and must not be interpreted as an observation-limited oracle.

4 DIAGNOSE

5 TRAINING

6 TEST AND TRANSFER

7 DISCUSSION

8 CONCLUSION

AI USE STATEMENT

(This section is **required** and does not count toward the page limit.)

In this work, we used generative AI tools for [tasks with required disclosure]. We have not used generative AI tools for [other tasks with required disclosure], and [the rest of the required disclosure tasks] are not applicable to this work. Additionally, we used generative AI tools for [tasks with recommended disclosure]. We have reviewed all AI-assisted work. [Elaborate. For example, “we checked LLM-generated research ideas for potential plagiarism through a manual literature survey”, “LLM-generated code was verified and tested for correctness by 2 authors”, etc.]. We take responsibility for the final content of this work, including text, claims or artifacts produced with the aid of generative AI.

See the ICLR 2027 AI Policy for Authors for more details. This statement should not be more than 1 page.

ETHICS STATEMENT

(This section is **recommended** and does not count toward the page limit.)

If authors feel that their paper submission raises questions regarding the Code of Ethics, they are encouraged to include a paragraph of Ethics Statement (at the end of the main text before references) to address potential concerns where appropriate. Topics include, but are not limited to, studies that involve human subjects, practices to data set releases, potentially harmful insights, methodologies and

applications, potential conflicts of interest and sponsorship, discrimination/bias/fairness concerns, privacy and security issues, legal compliance, and research integrity issues (e.g., IRB, documentation, research ethics). This statement should not be more than 1 page.

REPRODUCIBILITY STATEMENT

(This section is **recommended** and does not count toward the page limit.)

It is important that the work published in ICLR is reproducible. Authors are strongly encouraged to include a paragraph-long Reproducibility Statement at the end of the main text (before references) to discuss the efforts that have been made to ensure reproducibility. This paragraph should not itself describe details needed for reproducing the results, but rather reference the parts of the main paper, appendix, and supplemental materials that will help with reproducibility. For example, for novel models or algorithms, a link to an anonymous downloadable source code can be submitted as supplementary materials; for theoretical results, clear explanations of any assumptions and a complete proof of the claims can be included in the appendix; for any datasets used in the experiments, a complete description of the data processing steps can be provided in the supplementary materials. Each of the above are examples of things that can be referenced in the reproducibility statement.

AUTHOR CONTRIBUTIONS

If you'd like to, you may include a section for author contributions as is done in many journals. This is optional and at the discretion of the authors.

ACKNOWLEDGMENTS

Use unnumbered third level headings for the acknowledgments. All acknowledgments, including those to funding agencies, go at the end of the paper.

REFERENCES

- Leo Benac, Jonas Raedler, Zilin Ma, and Finale Doshi-Velez. A benchmark for multi-party negotiation games from real negotiation data. *arXiv preprint arXiv:2603.14066*, 2026. URL <https://arxiv.org/abs/2603.14066>.
- David Carmel and Shaul Markovitch. Opponent modeling in multi-agent systems. In *Adaptation and Learning in Multi-Agent Systems*, volume 1042 of *Lecture Notes in Computer Science*, pp. 40–52. Springer, 1996. doi: 10.1007/3-540-60923-7_18.
- Piotr J. Gmytrasiewicz and Prashant Doshi. A framework for sequential planning in multi-agent settings. *Journal of Artificial Intelligence Research*, 24:49–79, 2005. doi: 10.1613/JAIR.1579. URL <https://www.cs.cmu.edu/afs/cs/project/jair/pub/volume24/gmytrasiewicz05a.pdf>.
- Yizhe Huang, Anji Liu, Fanqi Kong, Yaodong Yang, Song-Chun Zhu, and Xue Feng. Efficient adaptation in mixed-motive environments via hierarchical opponent modeling and planning. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pp. 20004–20022. PMLR, 2024. URL <https://proceedings.mlr.press/v235/huang24p.html>.

A APPENDIX

B SUBMISSION OF CONFERENCE PAPERS TO ICLR 2027

ICLR requires electronic submissions, processed by <https://openreview.net/>. See ICLR's website for more instructions.

If your paper is ultimately accepted, the statement `\iclrfinalcopy` should be inserted to adjust the format to the camera ready requirements.

The format for the submissions is a variant of the NeurIPS format. Please read carefully the instructions below, and follow them faithfully.

B.1 STYLE

Papers to be submitted to ICLR 2027 must be prepared according to the instructions presented here.

Authors are required to use the ICLR L^AT_EX style files obtainable at the ICLR website. Please make sure you use the current files and not previous versions. Tweaking the style files may be grounds for rejection.

B.2 RETRIEVAL OF STYLE FILES

The style files for ICLR and other conference information are available online at:

<http://www.iclr.cc/>

The file `iclr2027_conference.pdf` contains these instructions and illustrates the various formatting requirements your ICLR paper must satisfy. Submissions must be made using L^AT_EX and the style files `iclr2027_conference.sty` and `iclr2027_conference.bst` (to be used with L^AT_EX2e). The file `iclr2027_conference.tex` may be used as a “shell” for writing your paper. All you have to do is replace the author, title, abstract, and text of the paper with your own.

The formatting instructions contained in these style files are summarized in sections C, D, and E below.

C GENERAL FORMATTING INSTRUCTIONS

The text must be confined within a rectangle 5.5 inches (33 picas) wide and 9 inches (54 picas) long. The left margin is 1.5 inch (9 picas). Use 10 point type with a vertical spacing of 11 points. Times New Roman is the preferred typeface throughout. Paragraphs are separated by 1/2 line space, with no indentation.

Paper title is 17 point, in small caps and left-aligned. All pages should start at 1 inch (6 picas) from the top of the page.

Authors’ names are set in boldface, and each name is placed above its corresponding address. The lead author’s name is to be listed first, and the co-authors’ names are set to follow. Authors sharing the same address can be on the same line.

Please pay special attention to the instructions in section E regarding figures, tables, acknowledgments, and references.

There will be a strict upper limit of **9 pages** for the main text of the initial submission, with unlimited additional pages for citations. This limit will be expanded to **10 pages** for rebuttal/camera ready.

D HEADINGS: FIRST LEVEL

First level headings are in small caps, flush left and in point size 12. One line space before the first level heading and 1/2 line space after the first level heading.

D.1 HEADINGS: SECOND LEVEL

Second level headings are in small caps, flush left and in point size 10. One line space before the second level heading and 1/2 line space after the second level heading.

D.1.1 HEADINGS: THIRD LEVEL

Third level headings are in small caps, flush left and in point size 10. One line space before the third level heading and 1/2 line space after the third level heading.

E CITATIONS, FIGURES, TABLES, REFERENCES

These instructions apply to everyone, regardless of the formatter being used.

E.1 CITATIONS WITHIN THE TEXT

Citations within the text should be based on the `natbib` package and include the authors' last names and year (with the "et al." construct for more than two authors). When the authors or the publication are included in the sentence, the citation should not be in parenthesis using `\citet{}` (as in "See ? for more information."). Otherwise, the citation should be in parenthesis using `\citep{}` (as in "Deep learning shows promise to make progress towards AI (?).").

The corresponding references are to be listed in alphabetical order of authors, in the REFERENCES section. As to the format of the references themselves, any style is acceptable as long as it is used consistently.

E.2 FOOTNOTES

Indicate footnotes with a number¹ in the text. Place the footnotes at the bottom of the page on which they appear. Precede the footnote with a horizontal rule of 2 inches (12 picas).²

E.3 FIGURES

All artwork must be neat, clean, and legible. Lines should be dark enough for purposes of reproduction; art work should not be hand-drawn. The figure number and caption always appear after the figure. Place one line space before the figure caption, and one line space after the figure. The figure caption is lower case (except for first word and proper nouns); figures are numbered consecutively.

Make sure the figure caption does not get separated from the figure. Leave sufficient space to avoid splitting the figure and figure caption.

You may use color figures. However, it is best for the figure captions and the paper body to make sense if the paper is printed either in black/white or in color.




Figure 1: Sample figure caption.

E.4 TABLES

All tables must be centered, neat, clean and legible. Do not use hand-drawn tables. The table number and title always appear before the table. See Table 1.

Place one line space before the table title, one line space after the table title, and one line space after the table. The table title must be lower case (except for first word and proper nouns); tables are numbered consecutively.

¹Sample of the first footnote

²Sample of the second footnote

Table 1: Sample table title

PART	DESCRIPTION
Dendrite	Input terminal
Axon	Output terminal
Soma	Cell body (contains cell nucleus)

F DEFAULT NOTATION

In an attempt to encourage standardized notation, we have included the notation file from the textbook, *Deep Learning* ? available at https://github.com/goodfeli/dlbook_notation/. Use of this style is not required and can be disabled by commenting out `math_commands.tex`.

Numbers and Arrays

a	A scalar (integer or real)
$\mathbf{a}$	A vector
$\mathbf{A}$	A matrix
$\mathbf{A}$	A tensor
I_n	Identity matrix with n rows and n columns
I	Identity matrix with dimensionality implied by context
$\mathbf{e}^{(i)}$	Standard basis vector $[0, \dots, 0, 1, 0, \dots, 0]$ with a 1 at position i
$\text{diag}(\mathbf{a})$	A square, diagonal matrix with diagonal entries given by $\mathbf{a}$
$\mathbf{a}$	A scalar random variable
$\mathbf{a}$	A vector-valued random variable
$\mathbf{A}$	A matrix-valued random variable

Sets and Graphs

$\mathbb{A}$	A set
$\mathbb{R}$	The set of real numbers
$\{0, 1\}$	The set containing 0 and 1
$\{0, 1, \dots, n\}$	The set of all integers between 0 and n
$[a, b]$	The real interval including a and b
$(a, b]$	The real interval excluding a but including b
$\mathbb{A} \setminus \mathbb{B}$	Set subtraction, i.e., the set containing the elements of $\mathbb{A}$ that are not in $\mathbb{B}$
$\mathcal{G}$	A graph
$\text{Pa}_{\mathcal{G}}(\mathbf{x}_i)$	The parents of $\mathbf{x}_i$ in $\mathcal{G}$

Indexing

432	a_i	Element i of vector $\mathbf{a}$, with indexing starting at 1
433	$\mathbf{a}_{-i}$	All elements of vector $\mathbf{a}$ except for element i
434	$A_{i,j}$	Element i, j of matrix $\mathbf{A}$
435	$\mathbf{A}_{i,:}$	Row i of matrix $\mathbf{A}$
436	$\mathbf{A}_{:,i}$	Column i of matrix $\mathbf{A}$
437	$\mathbf{A}_{i,j,k}$	Element (i, j, k) of a 3-D tensor $\mathbf{A}$
438	$\mathbf{A}_{:,:,i}$	2-D slice of a 3-D tensor
439	$\mathbf{a}_i$	Element i of the random vector $\mathbf{a}$
440	Calculus	
441	$\frac{dy}{dx}$	Derivative of y with respect to x
442	$\frac{\partial y}{\partial x}$	Partial derivative of y with respect to x
443	$\nabla_{\mathbf{x}} y$	Gradient of y with respect to $\mathbf{x}$
444	$\nabla_{\mathbf{X}} y$	Matrix derivatives of y with respect to $\mathbf{X}$
445	$\nabla_{\mathbf{X}} y$	Tensor containing derivatives of y with respect to $\mathbf{X}$
446	$\frac{\partial f}{\partial \mathbf{x}}$	Jacobian matrix $\mathbf{J} \in \mathbb{R}^{m \times n}$ of $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$
447	$\nabla_{\mathbf{x}}^2 f(\mathbf{x})$ or $\mathbf{H}(f)(\mathbf{x})$	The Hessian matrix of f at input point $\mathbf{x}$
448	$\int f(\mathbf{x}) d\mathbf{x}$	Definite integral over the entire domain of $\mathbf{x}$
449	$\int_{\mathbb{S}} f(\mathbf{x}) d\mathbf{x}$	Definite integral with respect to $\mathbf{x}$ over the set $\mathbb{S}$
450	Probability and Information Theory	
451	$P(\mathbf{a})$	A probability distribution over a discrete variable
452	$p(\mathbf{a})$	A probability distribution over a continuous variable, or over a variable whose type has not been specified
453	$\mathbf{a} \sim P$	Random variable $\mathbf{a}$ has distribution P
454	$\mathbb{E}_{\mathbf{x} \sim P}[f(\mathbf{x})]$ or $\mathbb{E}f(\mathbf{x})$	Expectation of $f(\mathbf{x})$ with respect to $P(\mathbf{x})$
455	$\text{Var}(f(\mathbf{x}))$	Variance of $f(\mathbf{x})$ under $P(\mathbf{x})$
456	$\text{Cov}(f(\mathbf{x}), g(\mathbf{x}))$	Covariance of $f(\mathbf{x})$ and $g(\mathbf{x})$ under $P(\mathbf{x})$
457	$H(\mathbf{x})$	Shannon entropy of the random variable $\mathbf{x}$
458	$D_{\text{KL}}(P \ Q)$	Kullback-Leibler divergence of P and Q
459	$\mathcal{N}(\mathbf{x}; \boldsymbol{\mu}, \boldsymbol{\Sigma})$	Gaussian distribution over $\mathbf{x}$ with mean $\boldsymbol{\mu}$ and covariance $\boldsymbol{\Sigma}$
460	Functions	
461		
462		
463		
464		
465		
466		
467		
468		
469		
470		
471		
472		
473		
474		
475		
476		
477		
478		
479		
480		
481		
482		
483		
484		
485		

486	$f : \mathbb{A} \rightarrow \mathbb{B}$	The function f with domain $\mathbb{A}$ and range $\mathbb{B}$
487	$f \circ g$	Composition of the functions f and g
488	$f(x; \theta)$	A function of x parametrized by θ . (Sometimes we write
489		$f(x)$ and omit the argument θ to lighten notation)
490		
491	$\log x$	Natural logarithm of x
492	$\sigma(x)$	Logistic sigmoid, $\frac{1}{1 + \exp(-x)}$
493	$\zeta(x)$	Softplus, $\log(1 + \exp(x))$
494	$\ \mathbf{x}\ _p$	L^p norm of $\mathbf{x}$
495	$\ \mathbf{x}\ $	L^2 norm of $\mathbf{x}$
496	x^+	Positive part of x , i.e., $\max(0, x)$
497	$\mathbf{1}_{\text{condition}}$	is 1 if the condition is true, 0 otherwise
498		
499		
500		
501		

502 G FINAL INSTRUCTIONS

503 Do not change any aspects of the formatting parameters in the style files. In particular, do not modify
 504 the width or length of the rectangle the text should fit into, and do not change font sizes (except
 505 perhaps in the REFERENCES section; see below). Please note that pages should be numbered.

506 H PREPARING POSTSCRIPT OR PDF FILES

507 Please prepare PostScript or PDF files with paper size “US Letter”, and not, for example, “A4”. The
 508 `-t letter` option on `dvips` will produce US Letter files.

509 Consider directly generating PDF files using `pdflatex` (especially if you are a MiKTeX user).
 510 PDF figures must be substituted for EPS figures, however.

511 Otherwise, please generate your PostScript and PDF files with the following commands:

```
512 dvips mypaper.dvi -t letter -Ppdf -G0 -o mypaper.ps
513 ps2pdf mypaper.ps mypaper.pdf
```

514 H.1 MARGINS IN LATEX

515 Most of the margin problems come from figures positioned by hand using `\special` or other
 516 commands. We suggest using the command `\includegraphics` from the `graphicx` package.
 517 Always specify the figure width as a multiple of the line width as in the example below using `.eps`
 518 graphics

```
519 \usepackage[dvips]{graphicx} ...
520 \includegraphics[width=0.8\linewidth]{myfile.eps}
```

521 or

```
522 \usepackage[pdftex]{graphicx} ...
523 \includegraphics[width=0.8\linewidth]{myfile.pdf}
```

524 for `.pdf` graphics. See section 4.4 in the graphics bundle documentation (<http://www.ctan.org/tex-archive/macros/latex/required/graphics/grfguide.ps>)

525 A number of width problems arise when LaTeX cannot properly hyphenate a line. Please give
 526 LaTeX hyphenation hints using the `\-` command.