

FORMATTING INSTRUCTIONS FOR ICLR 2027

CONFERENCE SUBMISSIONS

Anonymous authors

Paper under double-blind review

ABSTRACT

1 INTRODUCTION

Large language models (LLMs) are increasingly deployed as interacting agents that exchange information, negotiate decisions, and coordinate their actions over multiple turns. Such multi-agent systems (MAS) promise to combine information and capabilities distributed across multiple agents, but recent work has exposed substantial failures at the interaction level. [TODO: it is not too proper to mention decisive here cause we do not investigate this, better change to cannot actively acquire info, correctly interpret info from others. ground plans is also something we do not address.] Agents can overlook decisive information held by others, exchange extensive messages without reaching effective coordination, and fail to maintain grounded plans over multi-turn negotiation (Li et al., 2026a; Zhang et al., 2026b; Yao et al., 2026). [TODO: this sentence has low information, I believe these 2 cited paper here gives something more interesting to mention] Broader analyses similarly find that failures arising from inter-agent interaction constitute a distinct source of error beyond mistakes made by individual agents in isolation (Cemri et al., 2025; Anthropic, 2026). Given the growing use of agentic systems, understanding why individually capable agents can still coordinate poorly is an important problem.

[TODO: this is the place where we need to change the framing, as we discussed, we need to let the reviewer and readers know the case we are studying is important. also the logic connecting previous paragraph is not smooth. we said there are some issues in the previous paragraph, we should say how we are going to tackle not directly] We focus on a particularly consequential class of interactions: **coordination under partial information and mixed incentives**. These conditions arise whenever agents act on behalf of different participants who do not share all of their information or objectives. Scheduling agents, for example, coordinate using privately held constraints (Zou et al., 2026); agents in multi-party negotiation may cooperate while pursuing asymmetric private objectives (O’Neill et al., 2026); and autonomous agents representing separate principals can repeatedly interact while retaining distinct economic interests (Li et al., 2026b). These settings differ from fully cooperative systems with a common objective: successful interaction requires reasoning not only about the task, but also about other decision makers whose relevant information and incentives are only partially observable.

In this work, we examine one interaction-level difficulty that may contribute to failures in such settings. An agent must infer behaviorally relevant latent properties of other agents—such as what they know, prefer, or are likely to do—and use this information when choosing its own interaction. However, inference and decision making are not independent stages. What an agent currently infers changes the value of possible requests, disclosures, offers, or commitments; conversely, the interaction it chooses determines how other agents respond and therefore what evidence becomes available for subsequent inference. Thus, in the settings we study, reasoning repeatedly passes through the dependence

latent-state inference $\rightarrow$ interaction decision $\rightarrow$ new evidence $\rightarrow$ updated inference. (1)

This structure is closely related to classical models of interactive decision making in which uncertainty over other agents enters sequential planning (Gmytrasiewicz & Doshi, 2005; Huang et al., 2024). We ask whether current LLM agents can reliably perform these two computations together, and whether directly supervising them improves subsequent interaction.

Studying this question is difficult with end-task performance alone. Existing benchmarks provide evidence of failures in both inference and downstream coordination: agents fail to recover information distributed across participants (Li et al., 2026a), struggle with reasoning about other agents in coordination tasks (Agashe et al., 2025), and fail to maintain grounding and execute joint plans during negotiation (Yao et al., 2026). Yet the same unsuccessful outcome can arise for different reasons. An agent may choose a poor interaction because its inference about another agent is unsupported by the available evidence, or it may have an adequate inference but fail to act appropriately on it. To distinguish these cases, we need interaction states for which both the latent-state inference and the values of alternative actions can be evaluated.

Games provide a convenient substrate for constructing such states. Their information structure, latent variables, legal actions, and utilities can be explicitly controlled, and recent work shows that game interaction and self-play can train strategic capabilities that generalize beyond the training games (Yuan et al., 2026; Liu et al., 2026; Lyu et al., 2026; He et al., 2026b). Our use of games is motivated by a different requirement: **precise supervision of interaction reasoning**. Rather than relying exclusively on complete-game outcomes, we construct compact strategic slices of a game around decision points where an agent must infer relevant hidden state and act on that inference. This lets us specify not only whether the eventual game was won, but what could be inferred at the decision point and how different interactions should be valued.

To make this supervision possible, we extend commitment-based multi-party negotiation (Benac et al., 2026) with partial information and private player preferences. Players privately value public goals and interact through consequential offers, responses, and binding commitments. Their actions therefore simultaneously alter future strategic opportunities and provide evidence about latent preferences. At the same time, the game remains sufficiently compact to permit exact filtering over latent states, oracle partner policies, and counterfactual evaluation of candidate actions. We use this structure to construct **strategic interaction slices** that retain the information, incentives, and continuation required for meaningful multi-agent reasoning while remaining exactly analyzable.

We first use these slices to characterize failures in current LLM agents. Because the underlying latent states and continuation values are available to an oracle, we can separately measure errors in what the model infers from the interaction history and errors in how it acts under that inference. We then use the same structure to provide detailed supervision for both computations during game training. Importantly, our aim is not to establish that training on shortened game trajectories is generally superior to full-game self-play. The slices serve as a controlled mechanism for supplying supervision at strategically meaningful interaction states. We test whether training with this supervision improves behavior on unseen strategic configurations and, more importantly, whether the resulting interaction capability transfers to external multi-agent settings with different information structures and incentives.

Overall, our work uses a controllable game environment to turn an otherwise entangled multi-agent coordination problem into an interaction process that can be diagnosed and supervised at the decision level. Our contributions are threefold:

- **Diagnosing inference and decision making in strategic interaction.** We operationalize the sequential dependence between latent-state inference and interaction decisions in partial-information, mixed-incentive coordination, and separately evaluate failures in the two computations.
- **Constructing verifiable supervision through game design.** We show how compact game mechanics with private preferences and information-revealing actions can provide exact latent-state targets and counterfactual action values for training social interaction reasoning.
- **Training and testing transferable interaction capabilities.** We construct strategically meaningful game slices, use their oracle structure for detailed supervision, and evaluate whether the resulting improvements generalize to unseen strategic configurations and external multi-agent environments.

2 RELATED WORKS

Classical multi-agent decision making couples inference about other agents with downstream action selection, from opponent models and interactive beliefs (Carmel & Markovitch, 1996; Gmy-

trasiewicz & Doshi, 2005; Albrecht & Ramamoorthy, 2013) to belief-space coordination, inference from actions, and opponent-aware planning (Macindoe et al., 2012; Wu et al., 2021; Foerster et al., 2019; Huang et al., 2024). For LLM agents, explicit mental-state or behavioral representations support collaboration, adaptation, trust tracking, and planning (Li et al., 2023; Cross et al., 2025; Jha et al., 2026; Mu et al., 2026; Zhou et al., 2026; Feng et al., 2026; Goel et al., 2026; Hwang et al., 2026). However, accurate mental-state prediction need not yield useful behavior (Riemer et al., 2025; Sobotka et al., 2026).

Recent MAS benchmarks expose strategic failures under private information, mixed incentives, and imperfect communication. Commitment-based negotiation, self-negotiated contracts, and multi-party conquest model consequential agreements and coalition dynamics (Benac et al., 2026; Wyse et al., 2026; O’Neill et al., 2026); CalBench, SOTOPIA-TOM, and SNEAK make information disclosure and privacy strategically consequential (Zou et al., 2026; YS et al., 2026; Avsian & Heck, 2026). Broader evaluations cover collaboration, controlled interaction, long-horizon commerce, and communication barriers (Zhu et al., 2025; Chen et al., 2026; Li et al., 2026b; Xuan et al., 2026). Diagnostic work further finds failures in distributed information pooling, distributed computation, dynamic grounding, and belief-sensitive negotiation (Li et al., 2026a; Zhang et al., 2026b; Yao et al., 2026; Zhang et al., 2026a; Agashe et al., 2025); these complement broad failure taxonomies and accounts of long-lived peer interaction (Cemri et al., 2025; Anthropic, 2026).

Game-based training has been used to induce transferable strategic and social reasoning through self-play, multi-game reinforcement learning, and social reasoning supervision (Yuan et al., 2026; Liu et al., 2026; Lyu et al., 2026; He et al., 2026b; Hua et al., 2026; He et al., 2026a; Wu et al., 2026). Cross-environment cooperation and Other-Play study generalization across tasks and partners (Jha et al., 2025; Hu et al., 2020). Curriculum work selects terminal-proximal or high-learning-value subgames (West et al., 2020; Chen et al., 2024), while LLM curricula can also induce undesirable strategic abstractions (Madmoun & Lahlou, 2026). Similarity-based inference has also been proposed as a route to cooperation among foundation-model agents (Meulemans et al., 2026).

3 PRELIMINARIES

3.1 FROM INTERACTIVE UNCERTAINTY TO TWO CAPABILITIES

Coordination with private information is naturally an interactive partially observable decision problem. For agent i , an I-POMDP is

$$\mathcal{I}_i = \langle IS_i, A, T_i, \Omega_i, O_i, R_i \rangle, \quad IS_i = S \times \prod_{j \neq i} M_j, \quad (2)$$

where S is the physical state, M_j models agent j , A is the joint action space, and T_i, O_i, R_i specify transition, observation, and reward (Gmytrasiewicz & Doshi, 2005). Uncertainty therefore covers not only the world but also others’ information, preferences, and policies. This view underlies opponent modelling and belief-space coordination (Carmel & Markovitch, 1996; Huang et al., 2024; Goel et al., 2026); the observation–belief–action links nevertheless remain fragile for LLM agents (Sobotka et al., 2026; Jha et al., 2026).

Let $b_t^i \in \Delta(IS_i)$ be agent i ’s belief after public history h_t . An observation updates it as $b_{t+1}^i = \tau_i(b_t^i, a_t^i, o_{t+1}^i)$, while an action is selected by belief-conditioned planning,

$$a_t^{i,*} \in \arg \max_{a \in A_i} \mathbb{E} \left[\sum_{k=t}^T \gamma^{k-t} R_i(x_k, a_k) \mid b_t^i, a_t^i = a, \pi_{-i} \right]. \quad (3)$$

We call these computations *partner inference* (B), which maintains the part of b_t^i predictive of others, and *partner-conditioned planning* (P), which evaluates actions under that belief. They form the loop $B_t \rightarrow P_t \rightarrow a_t^i \rightarrow o_{t+1}^i \rightarrow B_{t+1}$: belief changes the best interaction, while the chosen interaction determines the next evidence. This is a functional decomposition, not a claim about separate neural modules.

3.2 BENAC-I: NEGOTIATION WITH PRIVATE PREFERENCES

To expose both computations while retaining exact supervision, we extend the commitment-based negotiation game BENAC (Benac et al., 2026). We call the resulting incomplete-information variant BENAC-I. An instance is

$$\mathcal{E} = \langle \mathcal{N}, \mathbf{K}, \mathcal{G}, \Gamma, \delta, \rho, C_0, \mathcal{L}, \{\Theta_i\}_{i \in \mathcal{N}}, \mu_0, \mathbf{q} \rangle. \quad (4)$$

Here $\mathcal{N}$ is the player set, K_i counts i 's binary commitment coordinates, and each public goal $g \in \mathcal{G}$ has requirements $\Gamma_g \subseteq \{(i, k) : k < K_i\}$ and mode $\delta_g \in \{\text{binary}, \text{linear}\}$. The proposer schedule ρ and legal-action relation $\mathcal{L}$ govern play from the initially empty, irreversible commitment state C_0 . Each player privately observes a preference type $v_i \in \Theta_i \subseteq \{-1, 0, +1\}^{|\mathcal{G}|}$, drawn from common prior μ_0 ; entries mean AVOID, NEUTRAL, or WANT. Vector $\mathbf{q}$ records remaining investigation budgets.

For commitment state C , goal satisfaction and terminal utility are

$$S_g(C) = \begin{cases} \mathbf{1}[c_{ik} = 1, \forall (i, k) \in \Gamma_g], & \delta_g = \text{binary}, \\ |\Gamma_g|^{-1} \sum_{(i, k) \in \Gamma_g} c_{ik}, & \delta_g = \text{linear}, \end{cases} \quad U_i(C; v_i) = \sum_g v_{ig} S_g(C). \quad (5)$$

On a proposal turn, a player may pass, offer new commitments (optionally as two alternatives), or use INVESTIGATE(j, g) when $q_i > 0$. Acceptance irreversibly binds both sides' proposed additions; rejection binds none; choosing from a two-offer proposal binds only the selected option. Investigation instead consumes the turn and one unit of q_i , leaves C unchanged, and privately reveals v_{jg} ; the target, but not the answer, is public. Negotiation therefore supplies indirect evidence and investigation supplies costly direct evidence within the same action space.

Reference teacher and labels. The teacher is an information-limited policy profile, not an omniscient player. At an information set $I_j = (h, v_j, r_j)$ —public history, own preferences, and private investigation results—it chooses

$$\text{supp } \pi_j^*(\cdot \mid I_j) \subseteq \arg \max_{a \in \mathcal{L}(I_j)} \mathbb{E}_{v_{-j} \sim b_j(\cdot \mid I_j)} [U_j(C_T; v_j) \mid a, \pi_{-j}^*], \quad (6)$$

breaking own-utility ties by total utility to others and mixing over residual ties. We enumerate the finite game tree, initialize legal actions uniformly, and synchronously update information-set best responses until the whole contingent profile is stable. We retain it only after a unilateral-deviation check; cycles and resource-limit failures yield no label. This defines a reproducible reference, not a unique or general-sum Nash equilibrium.

The same profile supplies both labels. An observed action gives the exact finite-support update $b_{t+1}^i(v_{-i}) \propto b_t^i(v_{-i}) \pi_j^*(u_t^j \mid I_t^j, v_j)$, treating i 's actions as interventions and conditioning on any private investigation result; this is the B target. For P , exact continuation gives $Q_i^*(I_i, a)$ for every legal action and hence every optimal action. Both labels thus share one declared behavioral model and terminal objective.

4 DIAGNOSE

We test whether a model can infer a partner's preference (B), act under a correct judgment (P), and compose the two. We freeze legal three-player histories and intervene on the judgment and first action, which end-game reward alone cannot do.

4.1 CONTROLLED NATIVE ENDGAMES

We evaluate Qwen3-4B-Instruct-2507 at temperature zero on 36 positions. Each preassigned stratum-split cell contains six distinct source games, with strata *known*, *unknown-irrelevant*, and *unknown-relevant*. Ego knows all but one partner-goal entry, whose possibilities are {WANT, NEUTRAL, AVOID}. We retain the public history, commitments, schedule, and all legal actions. Information-limited reference partners compute an exact terminal best response to a declared continuation. At most two ego decisions are measured, making counterfactual enumeration exact.

Table 1: Diagnosis on 12 unknown-relevant endgames. Exactness requires the full supported set; regret is in terminal utility units.

Quantity	Result
$B \rightarrow P$: type changes an optimal action	12/12
$P \rightarrow B$: action changes evidence entropy	12/12
Mean action loss / entropy contrast	1.50 / 0.94 bits
Root partner judgment exactly correct	1/12 (8.3%)
Planning regret with oracle judgment	0.56
Positive regret / exact post-evidence belief	8/12 / 26.4%

Selection uses game mechanics, never model answers. In an unknown-relevant position, an action optimal for one supported type is strictly suboptimal for another, certifying $B \rightarrow P$. The same position has two legal first actions whose response branches differ in expected posterior entropy, certifying $P \rightarrow B$. Unknown-irrelevant positions control for uncertainty that does not change the optimal-action set. These are task certificates, not claims about internal neural modules.

For B , the model returns every type compatible with history; unresolved sets are correct. For P , we supply the exact judgment and measure $\max_{a'} Q_i^*(h, b^*, a') - Q_i^*(h, b^*, a)$. We then force lower- and higher-information actions, enumerate all responses, and test the next belief, removing the model’s first choice as a confound.

4.2 BOTH CAPABILITIES FAIL WHERE BOTH ARE REQUIRED

Table 1 pools the two preassigned splits for the 12 unknown-relevant positions. In all 12, partner information changes both the best current action and what can be learned next.

The exact support is recovered once in 12 roots; the model averages 0.54 false exclusions. Even with the oracle judgment, regret is positive in eight cases (mean 0.56). After oracle-selected interactions, the next belief is exact on only 26.4% of probability-weighted branches. The model therefore both misreads evidence and fails to exploit a correct reading.

Representative failure patterns. Three cases make this concrete. (i) In native-30005-g0-d1, a partner’s offer is compatible with all three types, but the model treats a commitment used by the goal as a declaration of WANT and collapses the support. (ii) In the same state, even when given the correct support, it proposes a two-option offer to P_1 worth 0, while an offer to P_0 is worth 1: a planning error under an oracle judgment. (iii) In native-30040-g3-d2, an informative interaction separates a WANT branch from a {NEUTRAL, AVOID} branch. The model correctly outputs WANT after acceptance, yet also outputs it after the diagnostic rejection, reversing the evidence on that branch.

Known preferences do not eliminate planning error, while a wrong judgment need not matter when the entry is decision-irrelevant. Our scoped conclusion is that both B and P fail in states that couple them, and either failure can block a higher-value interaction. This motivates separate supervision of the two targets and explicit evaluation of their composition.

5 TRAINING

6 TEST AND TRANSFER

7 DISCUSSION

8 CONCLUSION

AI USE STATEMENT

We used generative AI tools to assist with literature discovery, reference formatting, and drafting and editing portions of this manuscript. All AI-assisted text and bibliographic metadata were reviewed

by the authors. The authors take responsibility for the final content of the paper, including its claims, experimental artifacts, and citations.

ETHICS STATEMENT

This work studies LLM-agent interaction in a synthetic, simulated negotiation environment and does not involve human participants or personal data. The environment includes mixed incentives and strategic information disclosure, which can surface behaviors such as manipulation or unhelpful information sharing. We use these behaviors only for controlled evaluation and training of coordination capabilities. We document the game rules, oracle policies, and limitations so that the resulting methods are not interpreted as deployment policies for real-world autonomous agents.

REPRODUCIBILITY STATEMENT

The main text and appendix specify the BENAC-P game rules, information structure, oracle belief filtering, and reference planning procedure. The appendix reports the diagnostic positions, prompts, evaluated model version, decoding configuration, controls, aggregation procedure, random seed, and provenance information used for the reported results. These details are intended to make the diagnostic suite and its measurements reproducible.

REFERENCES

- Saaket Agashe, Yue Fan, Anthony Reyna, and Xin Eric Wang. LLM-coordination: Evaluating and analyzing multi-agent coordination abilities in large language models. In *Findings of the Association for Computational Linguistics: NAACL 2025*, pp. 8053–8072. Association for Computational Linguistics, 2025. doi: 10.18653/v1/2025.findings-naacl.448.
- Stefano V. Albrecht and Subramanian Ramamoorthy. A game-theoretic model and best-response learning method for ad hoc coordination in multiagent systems. In *Proceedings of the 2013 International Conference on Autonomous Agents and Multi-Agent Systems*, pp. 1155–1156, 2013.
- Anthropic. Patterns and problems in emerging multiagent systems. Anthropic Research, August 2026. Published August 13, 2026.
- Adar Avsian and Larry Heck. SNEAK: Evaluating strategic communication and information leakage in large language models. In *Conference on Language Modeling (COLM)*, 2026.
- Leo Benac, Jonas Raedler, Zilin Ma, and Finale Doshi-Velez. A benchmark for multi-party negotiation games from real negotiation data, 2026.
- David Carmel and Shaul Markovitch. Opponent modeling in multi-agent systems. In *Adaptation and Learning in Multi-Agent Systems*, volume 1042 of *Lecture Notes in Computer Science*, pp. 40–52. Springer, 1996. doi: 10.1007/3-540-60923-7_18.
- Mert Cemri, Melissa Z. Pan, Shuyi Yang, Lakshya A. Agrawal, Bhavya Chopra, Rishabh Tiwari, Kurt Keutzer, Aditya Parameswaran, Dan Klein, Kannan Ramchandran, Matei Zaharia, Joseph E. Gonzalez, and Ion Stoica. Why do multi-agent LLM systems fail? In *Advances in Neural Information Processing Systems*, 2025.
- Jiaju Chen, Bo Sun, Yuxuan Lu, Yun Wang, Dakuo Wang, and Bingsheng Yao. CollabSim: A CSCW-grounded methodology for investigating collaborative competence of LLM agents through controlled multi-agent experiments, 2026.
- Jiayu Chen, Zelai Xu, Yunfei Li, Chao Yu, Jiaming Song, Huazhong Yang, Fei Fang, Yu Wang, and Yi Wu. Accelerate multi-agent reinforcement learning in zero-sum games with subgame curriculum learning. *Proceedings of the AAAI Conference on Artificial Intelligence*, 38(10):11320–11328, 2024. doi: 10.1609/aaai.v38i10.29011.
- Logan Cross, Violet Xiang, Agam Bhatia, Daniel L. K. Yamins, and Nick Haber. Hypothetical minds: Scaffolding theory of mind for multi-agent tasks with large language models. In *International Conference on Learning Representations*, 2025.

- Peilin Feng, Suorong Yang, and Soujanya Poria. Σ -Mem: An online reliability memory for LLM-based multi-agent systems, 2026.
- Jakob Foerster, Francis Song, Edward Hughes, Neil Burch, Iain Dunning, Shimon Whiteson, Matthew Botvinick, and Michael Bowling. Bayesian action decoder for deep multi-agent reinforcement learning. In *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pp. 1942–1951. PMLR, 2019.
- Piotr J. Gmytrasiewicz and Prashant Doshi. A framework for sequential planning in multi-agent settings. *Journal of Artificial Intelligence Research*, 24:49–79, 2005. doi: 10.1613/JAIR.1579.
- Harsh Goel, Aditya Sai Ellendula, Vaishnav Tadiparthi, Ehsan Moradi Pari, Hossein Nourkhiz Mahjoub, and Sandeep P. Chinchali. Bayesian partner modelling enables adaptive replanning for LLM coordination, 2026.
- Keyu He, Xuhui Zhou, and Maarten Sap. Social gym and SPaRTan: Benchmarking and improving LLM social reasoning via multi-agent game tournaments, 2026a.
- Yidong He, Yutao Lai, Pengxu Yang, Jiarui Gan, Jiexin Wang, Yi Cai, and Mengchen Zhao. Strat-reasoner: Reinforcing strategic reasoning of LLMs in multi-agent games. In *Proceedings of the 43rd International Conference on Machine Learning*, 2026b.
- Hengyuan Hu, Adam Lerer, Alex Peysakhovich, and Jakob Foerster. Other-Play for zero-shot coordination. In *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pp. 4399–4410. PMLR, 2020.
- Wenyue Hua, Zachary Huang, Tyler Payne, Safoora Yousefi, Saleema Amershi, and Asli Celikyilmaz. From passive delegates to strategic negotiators: Reinforcing social reasoning in small language models with SocialRL, 2026.
- Yizhe Huang, Anji Liu, Fanqi Kong, Yaodong Yang, Song-Chun Zhu, and Xue Feng. Efficient adaptation in mixed-motive environments via hierarchical opponent modeling and planning. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pp. 20004–20022. PMLR, 2024.
- EunJeong Hwang, Yuwei Yin, Giuseppe Carenini, Peter West, and Vered Shwartz. Infusing theory of mind into socially intelligent LLM agents. In *Findings of the Association for Computational Linguistics: ACL 2026*, pp. 11327–11360. Association for Computational Linguistics, 2026. doi: 10.18653/v1/2026.findings-acl.551.
- Kunal Jha, Wilka Carvalho, Yancheng Liang, Simon Shaolei Du, Max Kleiman-Weiner, and Natasha Jaques. Cross-environment cooperation enables zero-shot multi-agent coordination. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 27198–27220. PMLR, 2025.
- Kunal Jha, Aydan Huang, Eric Ye, Natasha Jaques, and Max Kleiman-Weiner. Modeling others’ minds as code. In *International Conference on Learning Representations*, 2026.
- Hua Li, Yu Quan Chong, Simon Stepputtis, Joseph Campbell, Dana Hughes, Michael Lewis, and Katia Sycara. Theory of mind for multi-agent collaboration via large language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pp. 180–192. Association for Computational Linguistics, 2023. doi: 10.18653/v1/2023.emnlp-main.13.
- Yuxuan Li, Aoi Naito, and Hirokazu Shirado. Systematic failures in collective reasoning under distributed information in multi-agent LLMs. In *Proceedings of the 43rd International Conference on Machine Learning*, volume 306 of *Proceedings of Machine Learning Research*, 2026a.
- Zeyuan Li, Lukas Petersson, Alessandro Acquisti, and Michiel A. Bakker. Emergent misaligned communication in long-horizon multi-agent LLM commerce, 2026b.

- Bo Liu, Leon Guertler, Simon Yu, Zichen Liu, Penghui Qi, Daniel Balcells, Mickel Liu, Cheston Tan, Weiyan Shi, Min Lin, Wee Sun Lee, and Natasha Jaques. SPIRAL: Self-play on zero-sum games incentivizes reasoning via multi-agent multi-turn reinforcement learning. In *International Conference on Learning Representations*, 2026.
- Nuoyan Lyu, Bingbing Xu, Weihao Meng, Yige Yuan, Yang Zhang, Zhiyong Huang, Tat-Seng Chua, and Huawei Shen. GIFT: Games as informal training for generalizable LLMs, 2026.
- Owen Macindoe, Leslie Pack Kaelbling, and Tomás Lozano-Pérez. POMCoP: Belief space planning for sidekicks in cooperative games. *Proceedings of the AAAI Conference on Artificial Intelligence and Interactive Digital Entertainment*, 8(1):38–43, 2012. doi: 10.1609/aiide.v8i1.12510.
- Hachem Madmoun and Salem Lahlou. Communication enables cooperation in LLM agents: A comparison with curriculum-based approaches. In *Proceedings of the 19th Conference of the European Chapter of the Association for Computational Linguistics (Volume 2: Short Papers)*, pp. 307–321. Association for Computational Linguistics, 2026. doi: 10.18653/v1/2026.eacl-short.23.
- Alexander Meulemans, Maciej Wołczyk, Marissa A. Weis, Rajai Nasser, Roberta Rocca, Seijin Kobayashi, Guillaume Lajoie, Angelika Steger, Blake Richards, Marcus Hutter, James Manyika, Rif A. Saurous, João Sacramento, and Blaise Agüera y Arcas. A game theory for foundation models shows new paths to rational cooperation through similarity inference, 2026.
- Chunjiang Mu, Ya Zeng, Qiaosheng Zhang, Kun Shao, Chen Chu, Hao Guo, Danyang Jia, Zhen Wang, and Shuyue Hu. Adaptive theory of mind for LLM-based multi-agent coordination. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, pp. 29608–29616, 2026. doi: 10.1609/aaai.v40i35.40204.
- Abigail O’Neill, Alan Zhu, Mihran Miroyan, Narges Norouzi, and Joseph E. Gonzalez. Cooperate to compete: Strategic coordination in multi-agent conquest. In *Proceedings of the 2026 Conference on Empirical Methods in Natural Language Processing*, 2026.
- Matthew Riemer, Zahra Ashktorab, Djallel Bouneffouf, Payel Das, Miao Liu, Justin D. Weisz, and Murray Campbell. Position: Theory of mind benchmarks are broken for large language models. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 82091–82130. PMLR, 2025.
- Jan Sobotka, Mustafa O. Karabag, and Ufuk Topcu. Why do LLMs struggle in strategic play? broken links between observations, beliefs, and actions, 2026.
- Joseph West, Frederic Maire, Cameron Browne, and Simon Denman. Improved reinforcement learning with curriculum. *Expert Systems with Applications*, 158:113515, 2020. doi: 10.1016/j.eswa.2020.113515.
- Jincenzi Wu, Yuxuan Lei, Jianxun Lian, Yitian Huang, Lexin Zhou, Haotian Li, Xing Xie, and Helen Meng. Social-R1: Towards human-like social reasoning in LLMs, 2026.
- Sarah A. Wu, Rose E. Wang, James A. Evans, Joshua B. Tenenbaum, David C. Parkes, and Max Kleiman-Weiner. Too many cooks: Bayesian inference for coordinating multi-agent collaboration. *Topics in Cognitive Science*, 13(2):414–432, 2021. doi: 10.1111/tops.12525.
- Tim Wyse, Kaitlin Bustos, Yulia Volkova, and Max Kleiman-Weiner. Commitment to cooperation with self-negotiated contracts. In *Proceedings of the Third Conference on Language Modeling (COLM)*, 2026.
- Keyang Xuan, Pengda Wang, Chongrui Ye, Haofei Yu, Tal August, and Jiaxuan You. SocialVeil: Probing social intelligence of language agents under communication barriers. In *Conference on Language Modeling (COLM)*, 2026.
- Yiheng Yao, Chelsea Zou, and Robert D. Hawkins. Talk is cheap, communication is hard: Dynamic grounding failures and repair in multi-agent negotiation, 2026. Accepted at the Conference on Language Modeling (COLM 2026).

- Yashwanth YS, Ruichen Wang, Shihua Zeng, Xuhui Zhou, Koichi Onoue, Vasudha Varadarajan, and Maarten Sap. SOTOPIA-TOM: Evaluating privacy and information management in multi-agent interaction with theory of mind. In *Conference on Language Modeling (COLM)*, 2026.
- Huining Yuan, Zelai Xu, Zheyue Tan, Xiangmin Yi, Mo Guang, Kaiwen Long, Haojia Hui, Boxun Li, Xinlei Chen, Bo Zhao, Xiao-Ping Zhang, Chao Yu, and Yu Wang. MARSHAL: Incentivizing multi-agent reasoning via self-play with strategic LLMs. In *International Conference on Learning Representations*, 2026.
- Erica Zhang, Fangzhao Zhang, Aneesh Pappu, Batu El, Jose Blanchet, Susan Athey, Jiashuo Liu, and James Zou. TERMS-Bench: Diagnosing LLM negotiation agents beyond deal rate, 2026a.
- Yuzhe Zhang, Feiran Liu, Yi Shan, Xinyi Huang, Xin Yang, Yueqi Zhu, Xuxin Cheng, Cao Liu, Ke Zeng, Terry Jingchen Zhang, and Wenyuan Jiang. SILO-BENCH: A scalable environment for evaluating distributed coordination in multi-agent LLM systems. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 29379–29398. Association for Computational Linguistics, 2026b. doi: 10.18653/v1/2026.acl-long.1354.
- Ruiwen Zhou, Maojia Song, Xiaobao Wu, Sitao Cheng, Xunjian Yin, Yuxi Xie, Zhuoqun Hao, Wenyue Hua, Liangming Pan, Soujanya Poria, and Min-Yen Kan. Epistemic context learning: Building trust the right way in LLM-based multi-agent systems, 2026.
- Kunlun Zhu, Hongyi Du, Zhaochen Hong, Xiaocheng Yang, Shuyi Guo, Zhe Wang, Zhenhailong Wang, Cheng Qian, Xiangru Tang, Heng Ji, and Jiaxuan You. MultiAgentBench: Evaluating the collaboration and competition of LLM agents. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 8580–8622. Association for Computational Linguistics, 2025. doi: 10.18653/v1/2025.acl-long.421.
- Chelsea Zou, Yiheng Yao, Selena She, Noah D. Goodman, and Robert D. Hawkins. CalBench: Evaluating coordination-privacy trade-offs in multi-agent LLMs, 2026.

A NATIVE DIAGNOSTIC PROTOCOL AND SUPPLEMENTARY RESULTS

Positions and oracle selection. Each frozen position contains its original game specification, independent per-player preference catalogues, ego’s preference row, a queried partner–goal pair and a replayable legal history. The ten positions have three players, two binary commitments per player, eight public goals and two scheduled proposal turns per player. At most three original proposal turns remain at the root. No preparation, assessment or automatic approval action is introduced. The active partner policy uses public history and its own preferences; its terminal best response is to the declared fixed continuation, not to an omniscient or equilibrium opponent.

For the judgment-sensitive cohort, the oracle evaluates all legal actions under each still-supported partner type. Admission requires an action optimal under one type with positive regret under another. This witnesses a consequential dependence on partner understanding, while allowing actions that are optimal under multiple types. It neither requires the model’s judgment to be wrong nor guarantees a change when that particular judgment is repaired. For the action–evidence cohort, native enumeration selects the first pair of legal actions with different expected posterior entropy and a nonterminal next ego decision on every branch. “Higher” and “lower” information refer only to this pair, not global information extrema. The utility-optimal reference action is selected independently. A positive downstream value from updating a frozen prior is not an admission condition.

A.1 WORKED MENU EXAMPLE

In `native-30040-g6-d2`, let x_i denote the already-bound $P_i.A_0$ and y_i denote $P_i.A_1$. Initially all x_i are true and all y_i are false. Table 3 gives the complete game; only v , P_2 ’s preference for G_6 , varies across the three compatible types. The public prefix is P_0 pass; a P_1 menu choosing between binding $(P_1.A_0, P_2.A_0)$ and $(P_1.A_1, P_2.A_0)$, with P_2 choosing the former; P_2 offering to add $P_0.A_0$, accepted by P_0 ; and P_2 offering to add $P_1.A_1$, rejected by ego. The resulting state is the example’s root.

Table 2: Frozen native positions. D and C retain the source-seed discovery and confirmation labels used during development; C is not an untouched test set. Each direction has three games in each partition.

Position identifier	Partition	$B \rightarrow P$	$P \rightarrow B$
native-30003-g6-d2	C	yes	
native-30005-g0-d1	C	yes	
native-30009-g2-d2	C	yes	
native-30013-g6-d3	C		yes
native-30016-g5-d1	D	yes	yes
native-30018-g5-d2	D	yes	yes
native-30019-g2-d1	C		yes
native-30020-g5-d1	D	yes	
native-30041-g6-d1	C		yes
native-30040-g6-d2	D		yes

Table 3: Complete goal and preference specification of the worked example. Entries +1, 0 and −1 mean WANT, NEUTRAL and AVOID. Here $v \in \{+1, 0, -1\}$ is hidden from ego.

Goal	Required commitments	P_0	P_1 (ego)	P_2
G_0	$x_0 \wedge y_1$	0	+1	0
G_1	$x_0 \wedge y_1 \wedge x_2$	−1	+1	+1
G_2	$y_0 \wedge x_1 \wedge y_2$	+1	−1	−1
G_3	$y_0 \wedge y_1 \wedge y_2$	0	+1	−1
G_4	$x_0 \wedge x_1$	+1	+1	0
G_5	$x_1 \wedge y_2$	0	−1	+1
G_6	$y_0 \wedge y_2$	+1	+1	v
G_7	$y_1 \wedge x_2$	−1	−1	+1

The menu offers O_1 , adding only y_2 , and O_2 , adding y_1 and y_2 . Exact terminal values for P_2 , evaluated under its declared fixed reference continuation, are

v	REJECT	CHOOSE_1	CHOOSE_2	oracle response
+1	1	1	2	CHOOSE_2
0	0	1	1	CHOOSE_2
−1	0	1	0	CHOOSE_1

For the neutral type, the immediate-reference tie-break selects O_2 : it immediately satisfies P_2 ’s wanted G_1, G_5, G_7 , whereas O_1 satisfies only G_5 . Thus CHOOSE_1 identifies AVOID, and CHOOSE_2 leaves WANT/NEUTRAL. Offering O_1 alone elicits acceptance from all three types. The comparison illustrates how the menu changes the evidence channel without changing the legal game protocol.

After either menu choice, the actual active P_0 offers to add y_0 while asking ego for no additional commitment. After O_1 , ego’s terminal values for reject and accept are both 0. After O_2 , they are 1 and 2 respectively. The difference follows from the bound commitments: adding y_0 after O_1 activates an avoided G_2 and a wanted G_6 , which cancel; after O_2 it also activates the wanted G_3 . The observed model judgment after the O_1 menu branch is {WANT, NEUTRAL}, excluding the only supported type. This example demonstrates the action–evidence contrast and an inference error; its between-action utility difference also reflects physical commitments. It is one of the six action–evidence cases, not an additional replication.

A.2 MEASUREMENT AND EXECUTION DETAILS

Queries and interventions. Let $B^*(s, H)$ denote the evidence-supported preference set and $\hat{B}(s, H)$ the model’s reported set. Both paired planning queries receive the same state s , complete public history H , initial possibilities and legal actions. Only the explicit judgment differs. The belief-generation scratchpad is not passed to either fresh planning context. For a supplied judgment B , write the model action as $\hat{a}(B)$ and let $Q^*(a)$ be its exact expected terminal value under the evidence-supported reference judgment. We report

$$\text{Reg}(a) = \max_{a'} Q^*(a') - Q^*(a), \quad \Delta_B = \text{Reg}(\hat{a}(\hat{B})) - \text{Reg}(\hat{a}(B^*)). \quad (7)$$

Positive Δ_B means beneficial repair; either sign can accompany an action change. Exact-set B accuracy requires every supported preference and no unsupported preference. Reference beliefs are uniform over compatible types; all probabilities and action values are computed in the backend.

For forced action a , let E_a be the subsequent sequence of observed partner actions before the next ego decision, and let Z be the uncertain partner type. The channel statistic is $\mathbb{E}[H(Z | H, a, E_a)]$ in bits, with every branch weighted by its compatible-type mass. The information gap subtracts this statistic for the higher-information arm from that for the lower-information arm. Model exactness is scored against the appropriate support on each branch, then averaged with the same weights. Terminal branches in other action arms require only a judgment; no fictitious continuation action is added. After the second measured ego action, the reference planner controls the remaining continuation.

Execution and semantic controls. We use Qwen3-4B-Instruct-2507 with temperature zero, served through vLLM 0.28 and the Hermes tool-call parser. The prompt requests brief reasoning with soft targets of 120 words for judgments and 240 for actions. Each response exposes one submission tool with automatic tool choice and a 1,024-token completion cap. Only truncation permits one additional 128-token submission request, without corrective feedback. All 127 formal tasks were valid on their first request; no finalization was needed. All contained assistant text before the tool call, averaging 131.4 whitespace-delimited words and 230.8 completion tokens per task (29,308 completion tokens in total). These traces are observable explanations, not a validated measure of the model’s internal computation.

Four separate empty-history controls supply singleton WANT, NEUTRAL, AVOID, or the unresolved three-label set. All four returned valid submissions, but the WANT control incorrectly returned all three labels. Its explanation incorrectly described the supplied initial set as containing all three values. The run initially stopped under the semantic gate and then continued under a protocol-validity gate, reusing all four cached responses, including this error. Neither prompts nor inference budgets were changed, and incorrect semantic answers were not retried. The controls consumed 491 completion tokens in four requests, separately from formal-task counts. This failure limits attribution of B errors: some can reflect reading or prior-grounding failures rather than strategic inference alone.

Aggregation and partition results. Panel A of Table 1 counts the six certified positions in each cohort. These counts summarize admission criteria, not six successful model interventions. The judgment-sensitive certificate uses two compatible type-conditioned singleton judgments at the same public state and history; it does not claim that the actual history identifies either singleton. The model’s oracle-judgment input instead contains the full evidence-supported set. Panel B reports model errors under the controlled inputs, separately from these task-level certificates.

The main table combines the two development partitions *within* each six-game cohort. It never pools the overlapping cohorts as twelve independent games. Branches are weighted within games before computing game means. The 95% intervals use 2,000 percentile-bootstrap resamples of source games with random seed zero. They describe variation in this selected small set, not failure prevalence across random games or an untouched confirmatory test. The 0.918-bit channel gap is identical in all six selected pairs; its absence of between-game variation is a property of these examples, not a precision claim about other environments.

Table 4: Results by retained development partition, three games per cell. Regret and judgment repair use the cohort roots. Forced-arm B accuracy is branch-weighted. Judgment replacement changes one action in C and two in D.

Measure	D	C
<i>Judgment-sensitive cohort</i>		
B exact	0.0%	33.3%
P regret with reference judgment	0.389	0.500
Judgment repair: reduction in regret	−0.222	0.000
<i>Action-evidence cohort</i>		
Root P regret with reference judgment	0.667	1.333
B exact, lower-information action	0.0%	33.3%
B exact, higher-information action	11.1%	33.3%

Task consequences and secondary decompositions. The following joint-response measurements supplement the oracle certificates and interface-specific error measurements; their sign is not an admission criterion or a test of whether the task dependency exists. Across the six judgment-sensitive cases, mean model-planner regret reduction is -0.111 $[-0.333, 0.000]$. Across the six action-evidence cases, the paired difference in B exactness is 5.6 $[-38.9, 44.4]$ percentage points. Neither quantity establishes a positive repair benefit.

In the judgment-sensitive cohort, replacing the explicit judgment changes `native-30005-g0-d1` from a menu to an offer and `native-30016-g5-d1` from a menu to pass, with no change in expected utility. At `native-30018-g5-d2`, it changes pass to an offer and increases regret from zero to $2/3$. These are observed effects of the supplied judgment on the same planner, not evidence that erroneous judgments are generally preferable. The other three root actions are unchanged.

For the forced action pairs, the higher-minus-lower reference utility contrast is 0.444 $[-0.222, 1.111]$; with the model choosing the second action, it is -0.222 $[-1.333, 0.833]$. These quantities include physical commitment changes and do not identify information-mediated utility effects. The auxiliary chooser/updater decomposition has ten complete tables, but replacing the updater changes reference-continuation utility by zero in both the utility-optimal and model-selected first-action arms in every game. We do not use these zeros as evidence that inference is irrelevant, or claim a positive updater-repair benefit. One root reference-planner score under a wrong model judgment is undefined because the assumed type contradicts the public history; that secondary score is marked unavailable, not filled with zero. Direct model action regret and the paired model-planner intervention remain defined.

Relation to the preceding native diagnosis. An earlier native run covered 36 positions from 26 source games, with known, unknown/relevant and unknown/irrelevant conditions. Its six designated confirmation unknown/relevant primary cases had B exact agreement 16.7% $[0.0, 50.0]\%$ and reference-judgment P regret 0.667 $[0.250, 1.083]$. Those primary cases mixed root and post-interaction decisions. That run produced 301 valid submissions out of 302 tasks and passed all four semantic controls. We retain it as preceding B/P evidence, without pooling it with this supplement: nine positions are reused, and the additional position is another query/history from an already represented source game. The earlier restricted three-stage protocol, which changed action structure and coupled partner types, is not the source of any results reported here.